

# FAKE NEWS DETECTION USING NLP

Technology: ARTIFICIAL INTELLIGENCE

NAME: MD SUHAIB V

NM ID: au311121205036

REG. NO.: 311121205036

## PHASE 4 - PROJECT SUBMISSION

### Phase 4: Development Part 2


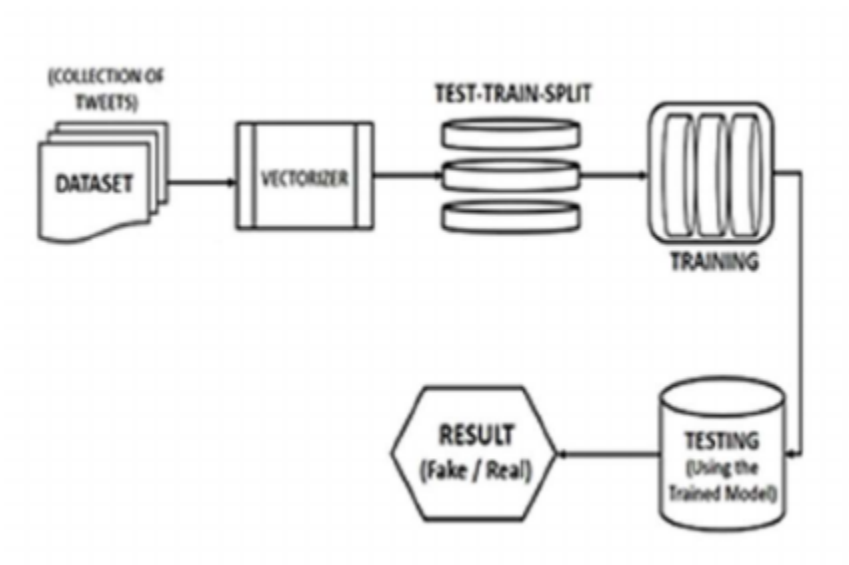


Figure 1. Generic LSTM architecture.

## **PHASE 3 GUIDELINES:(GIVEN)**

### **Phase 4: Development Part 2**

In this part you will continue building your project.

Continue building the fake news detection model by applying NLP techniques and training a classification model.

Text Preprocessing and Feature Extraction

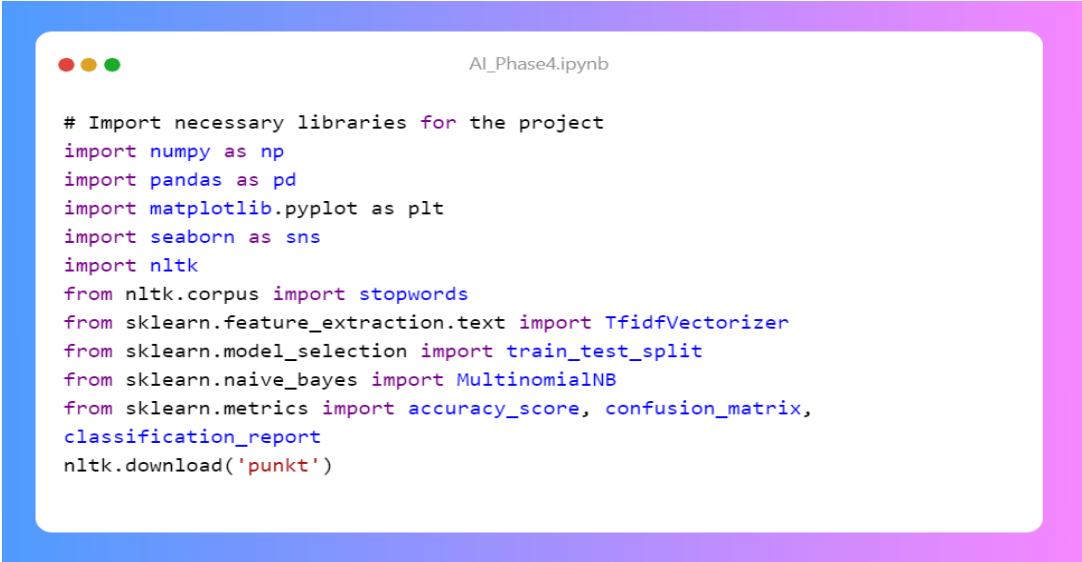
Model training and evaluation

## **INTRODUCTION:**

In phase 4 of this project, we continue with the further development of our fake news detection project. In the last phase of development we focused on loading the dataset and preprocessing the textual dataset. In this phase, we will proceed further with processes like: **Text Preprocessing, Feature Extraction, Applying NLP techniques and training a classification model, Model training and evaluation.**

Let us begin with the project first by seeing the steps we have already seen before moving on to the rest of the steps.

## **STEP 1: IMPORTING THE LIBRARIES.**




```
AI_Phase4.ipynb

# Import necessary libraries for the project
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import nltk
from nltk.corpus import stopwords
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
nltk.download('punkt')
```

Here, we import the necessary libraries and modules for the project. Let's break down their roles:

- numpy and pandas are fundamental for data manipulation and numerical operations.
- matplotlib and seaborn are used for data visualization, providing tools to create plots and graphs.
- nltk (Natural Language Toolkit) is a library for natural language processing (NLP) tasks, including text processing.
- nltk.corpus is used to access stopwords, which are common words to be excluded in text preprocessing.
- TfidfVectorizer from sklearn.feature\_extraction.text is a tool to convert text data into numerical features using TF-IDF.
- train\_test\_split from sklearn.model\_selection is used to split the data into training and testing sets.
- MultinomialNB from sklearn.naive\_bayes is a Naive Bayes classifier suitable for text classification.
- accuracy\_score, confusion\_matrix, and classification\_report from sklearn.metrics are for model evaluation.
- nltk.download('punkt'): This command downloads the 'punkt' resource for NLTK. The 'punkt' resource includes data files used for tokenization, which is the process of splitting text into individual words or tokens.

## **STEP 2: DATA LOADING:**



```
AI_Phase4.ipynb

# Load the dataset from Kaggle
fake_news = pd.read_csv("/Fake.csv")
real_news = pd.read_csv("/True.csv")
```

## True.csv

## Fake.csv

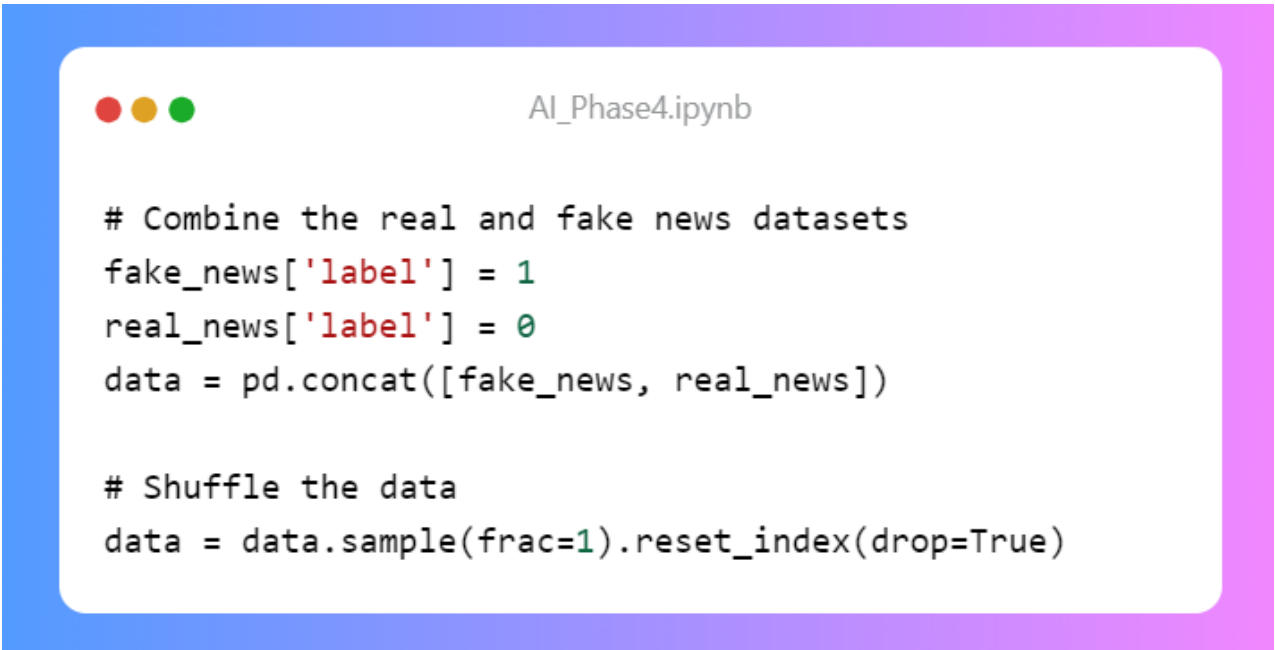
[illegible]

In this section, we load the dataset from Kaggle.

The datasets are downloaded from kaggle and loaded.

The `pd.read_csv()` function reads data from CSV files and stores it in Pandas DataFrames. The `fake_news` and `real_news` DataFrames will contain the fake and real news data, respectively.

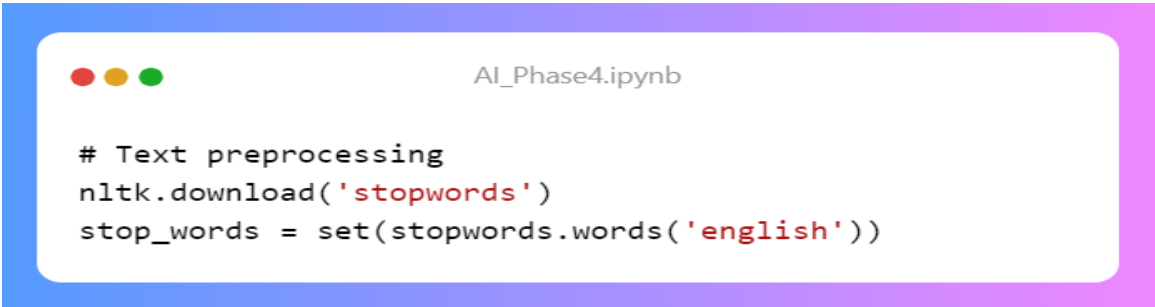
### **STEP 3: DATA PREPROCESSING:**

A screenshot of a Jupyter Notebook code cell. The cell has a title bar with three colored circles (red, yellow, green) and the text "AI\_Phase4.ipynb". The code inside the cell is as follows:

```
# Combine the real and fake news datasets
fake_news['label'] = 1
real_news['label'] = 0
data = pd.concat([fake_news, real_news])

# Shuffle the data
data = data.sample(frac=1).reset_index(drop=True)
```

In this section, we combine the fake and real news datasets into one DataFrame called `data`. We add a 'label' column, where '1' indicates fake news and '0' indicates real news. The data is shuffled to ensure randomness, and the index is reset.

A screenshot of a Jupyter Notebook code cell. The cell has a title bar with three colored circles (red, yellow, green) and the text "AI\_Phase4.ipynb". The code inside the cell is as follows:

```
# Text preprocessing
nltk.download('stopwords')
stop_words = set(stopwords.words('english'))
```

Here, we download the list of English stopwords from NLTK, which are common words that don't carry significant meaning in text data.

```
AI_Phase4.ipynb

def preprocess_text(text):
    # Lowercase the text
    text = text.lower()
    # Tokenize the text
    tokens = nltk.word_tokenize(text)
    # Remove stopwords
    tokens = [word for word in tokens if word not in stop_words]
    # Join the tokens back into a string
    text = ' '.join(tokens)
    return text

data['text'] = data['text'].apply(preprocess_text)
```

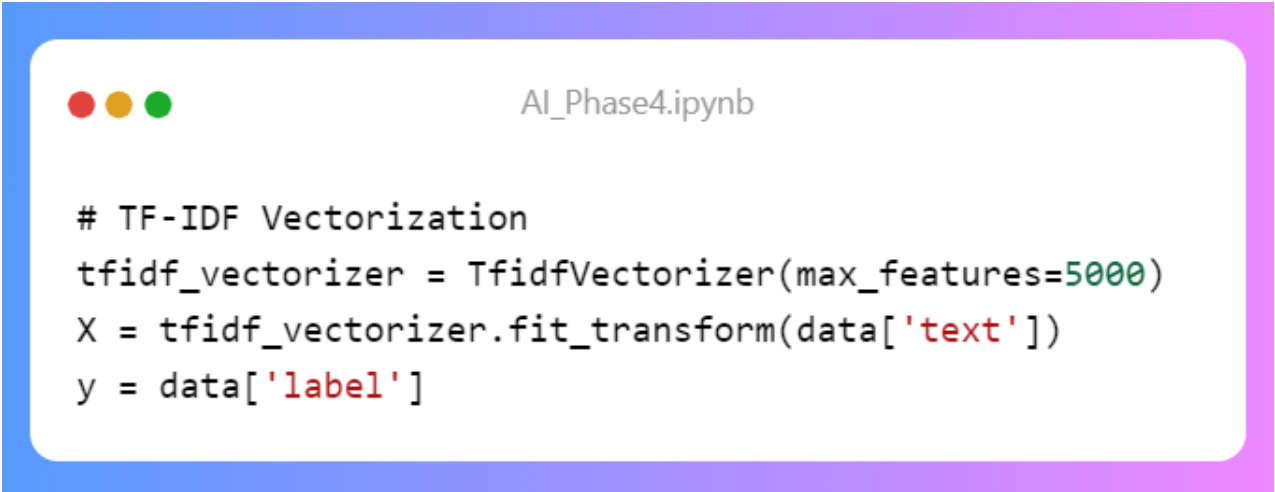
This part defines a function `preprocess_text` to apply text preprocessing to the 'text' column in the DataFrame. The steps are as follows:

- Convert the text to lowercase to make it uniform.
- Tokenize the text into individual words using NLTK.
- Remove common English stopwords from the tokenized words.
- Join the tokenized words back into a cleaned text.
- 

```
[8] print(data['text'])

0      claiming least racist person , donald trump li...
1      dunkin donuts american global donut company co...
2      london ( reuters ) - britain made substantive ...
3      washington ( reuters ) - republican party ' tw...
4      anyone else wondering cop-hating , racist , be...
...
44893   30 years , donald trump built great negotiator...
44894   washington ( reuters ) - united states monday ...
44895   sarah palin marked return national stage part ...
44896   seoul ( reuters ) - south korean president moo...
44897   samarkand , uzbekistan ( reuters ) - senior of...
Name: text, Length: 44898, dtype: object
```

## **STEP 4: FEATURE EXTRACTION:**



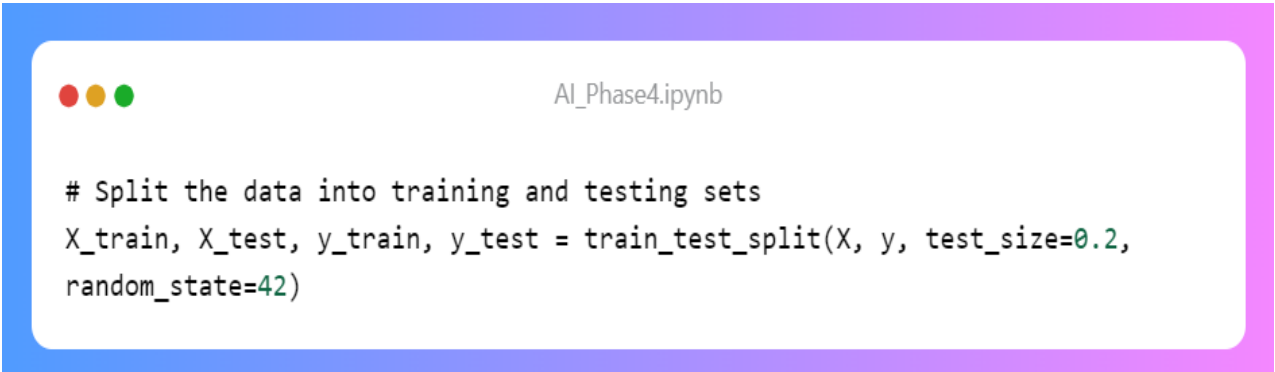
AI\_Phase4.ipynb

```
# TF-IDF Vectorization
tfidf_vectorizer = TfidfVectorizer(max_features=5000)
X = tfidf_vectorizer.fit_transform(data['text'])
y = data['label']
```

In this section, we use TF-IDF vectorization to convert the preprocessed text data into numerical features. Let's explain this step by step:

- TfidfVectorizer is initialized with max\_features set to 5000. This specifies that we want to consider the top 5000 most important terms in the dataset.
- tfidf\_vectorizer.fit\_transform(data['text']) computes the TF-IDF scores for each term in the text data, and X is a sparse matrix that represents the transformed data.
- y contains the labels for the corresponding data.

## **STEP 5: MODEL TRAINING AND EVALUATION:**

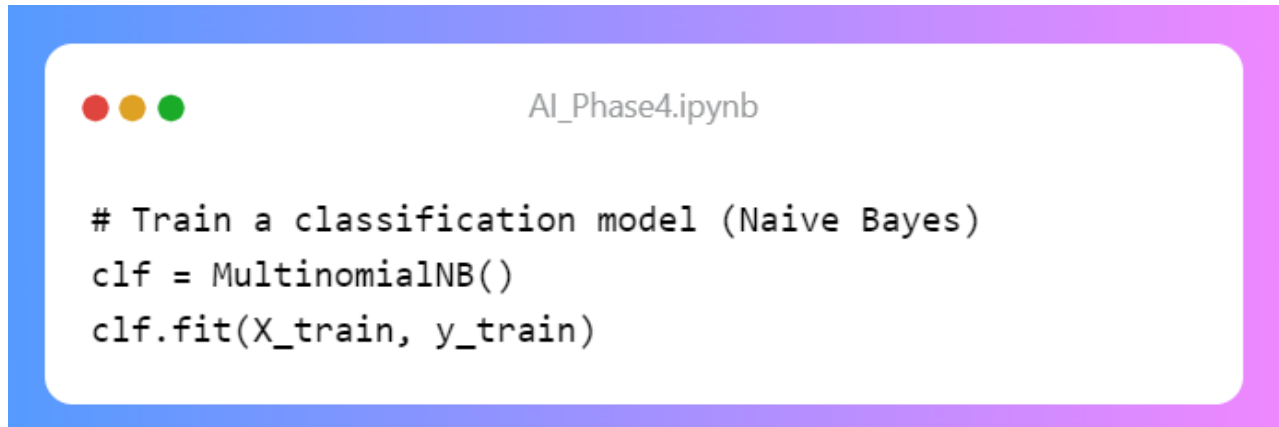


AI\_Phase4.ipynb

```
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

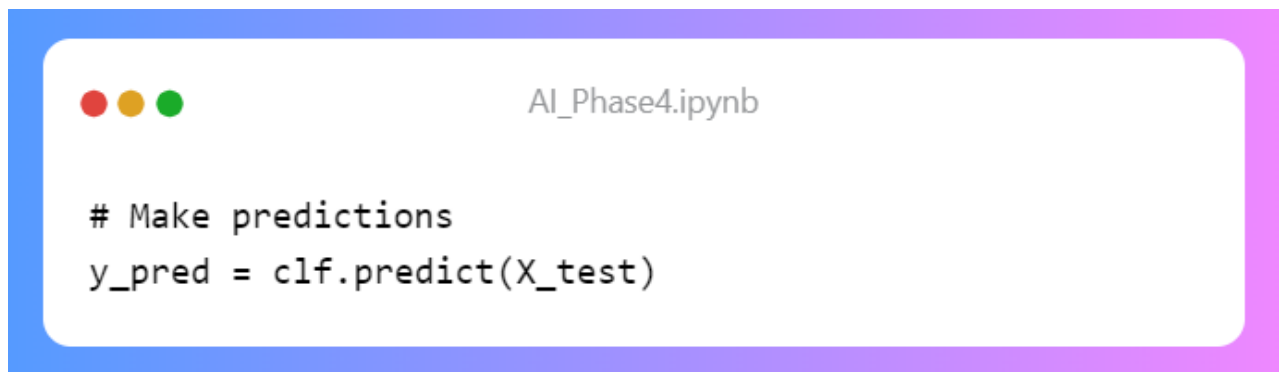
In this part, the data is split into training and testing sets. Here's what each line does:

- `train_test_split(X, y, test_size=0.2, random_state=42)` splits the feature matrix `X` and the labels `y` into training and testing sets.
- The `test_size` parameter specifies the percentage of data to use for testing (20% in this case), and `random_state` ensures reproducibility.

A screenshot of a Jupyter Notebook code cell. The cell has a title bar with three colored circles (red, yellow, green) and the text "AI\_Phase4.ipynb". The code inside the cell is as follows:

```
# Train a classification model (Naive Bayes)
clf = MultinomialNB()
clf.fit(X_train, y_train)
```

Here, we create a Multinomial Naive Bayes classifier **clf** and train it using the training data. Naive Bayes is a simple yet effective classification algorithm, commonly used in text classification tasks.

A screenshot of a Jupyter Notebook code cell. The cell has a title bar with three colored circles (red, yellow, green) and the text "AI\_Phase4.ipynb". The code inside the cell is as follows:

```
# Make predictions
y_pred = clf.predict(X_test)
```

The model is used to make predictions on the testing set, and the predictions are stored in the `y_pred` variable.





AI\_Phase4.ipynb

```
# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
confusion = confusion_matrix(y_test, y_pred)
classification_rep = classification_report(y_test, y_pred)
```

Here, we evaluate the model's performance.

- `accuracy_score` calculates the accuracy of the model's predictions.
- `confusion_matrix` generates a confusion matrix, which shows the number of true positives, true negatives, false positives, and false negatives.
- `classification_report` provides a comprehensive report including precision, recall, F1-score, and support for each class.
- `print("Accuracy:", accuracy)`, `print("Confusion Matrix:\n", confusion)`, and `print("Classification Report:\n", classification_rep)` display the evaluation results to the console.



AI\_Phase4.ipynb

```
print("Accuracy:", accuracy)
print("Confusion Matrix:\n", confusion)
print("Classification Report:\n", classification_rep)
```



Accuracy: 0.9359688195991092

Confusion Matrix:

[[3983 301]

[ 274 4422]]

Classification Report:

	precision	recall	f1-score	support
0	0.94	0.93	0.93	4284
1	0.94	0.94	0.94	4696
accuracy			0.94	8980
macro avg	0.94	0.94	0.94	8980
weighted avg	0.94	0.94	0.94	8980

## **EXAMPLE PROGRAM USING LOGISTIC REGRESSION AND NEURAL NETWORKS:**

```
import pandas as pd
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score,
roc_auc_score
from sklearn.linear_model import LogisticRegression
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense

# Load the "Fake.csv" dataset
fake_data = pd.read_csv("C:\\Users\\Bylee\\Downloads\\Fake.csv\\Fake.csv")

# Load the "True.csv" dataset
true_data = pd.read_csv("C:\\Users\\Bylee\\Downloads\\True.csv\\True.csv")
```

```
# Add labels to distinguish between fake and true news
fake_data['label'] = 0 # 0 for fake news
true_data['label'] = 1 # 1 for true news

# Combine the datasets
combined_data = pd.concat([fake_data, true_data], ignore_index=True)

# Data Preprocessing
combined_data['text'] = combined_data['title'] + " " + combined_data['text']

# Feature Extraction (TF-IDF)
tfidf_vectorizer = TfidfVectorizer(max_features=5000)
tfidf_matrix = tfidf_vectorizer.fit_transform(combined_data['text'])

# Model Selection
X_train, X_test, y_train, y_test = train_test_split(tfidf_matrix,
combined_data['label'], test_size=0.2, random_state=42)

# Logistic Regression Model
logistic_regression_model = LogisticRegression()
logistic_regression_model.fit(X_train, y_train)

# Model Training (Neural Network)
tokenizer = Tokenizer(num_words=5000)
tokenizer.fit_on_texts(combined_data['text'])
X_train_nn = tokenizer.texts_to_sequences(combined_data['text'])
X_train_nn = pad_sequences(X_train_nn, maxlen=100)

model = Sequential()
model.add(Embedding(input_dim=5000, output_dim=128, input_length=100))
model.add(LSTM(128))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
model.fit(X_train_nn, combined_data['label'], epochs=5, batch_size=64)
```

```
# Evaluation
```

```
# For Logistic Regression
```

```
y_pred = logistic_regression_model.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
precision = precision_score(y_test, y_pred)
```

```
recall = recall_score(y_test, y_pred)
```

```
f1 = f1_score(y_test, y_pred)
```

```
roc_auc = roc_auc_score(y_test, y_pred)
```

```
print(f"Logistic Regression Accuracy: {accuracy}")
```

```
print(f"Logistic Regression Precision: {precision}")
```

```
print(f"Logistic Regression Recall: {recall}")
```

```
print(f"Logistic Regression F1-Score: {f1}")
```

```
print(f"Logistic Regression ROC-AUC: {roc_auc}")
```

```
# For Neural Network
```

```
X_test_nn = tokenizer.texts_to_sequences(combined_data['text'])
```

```
X_test_nn = pad_sequences(X_test_nn, maxlen=100)
```

```
loss, accuracy = model.evaluate(X_test_nn, combined_data['label'])
```

```
print(f"Neural Network Accuracy: {accuracy}")
```

## **1. Importing Libraries:**

```
AI_Phase4.ipynb

import pandas as pd
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score,
f1_score, roc_auc_score
from sklearn.linear_model import LogisticRegression
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense
```

In this part, we import the necessary libraries and modules for the project. Here's what each library/module does:

- pandas is used for data manipulation.
- TfidfVectorizer from sklearn.feature\_extraction.text is used for TF-IDF feature extraction.
- train\_test\_split from sklearn.model\_selection splits the data into training and testing sets.
- accuracy\_score, precision\_score, recall\_score, f1\_score, and roc\_auc\_score from sklearn.metrics are used for model evaluation.
- LogisticRegression from sklearn.linear\_model is for training a logistic regression model.
- Tokenizer and pad\_sequences from tensorflow.keras.preprocessing.text are for text tokenization and padding sequences.
- Sequential, Embedding, LSTM, and Dense from tensorflow.keras.models and tensorflow.keras.layers are for building a neural network model.

## 2. Data Loading

A screenshot of a Jupyter Notebook interface. At the top, there are three colored circles (red, yellow, green) and the text "AI\_Phase4.ipynb". Below this, the code is displayed in a monospaced font. The code consists of three lines: a comment line "# Load the dataset from Kaggle", followed by "fake\_data = pd.read\_csv("/Fake.csv")", and "true\_data = pd.read\_csv("/True.csv")". The file names "Fake.csv" and "True.csv" are highlighted in red in the original image.

```
AI_Phase4.ipynb

# Load the dataset from Kaggle
fake_data = pd.read_csv("/Fake.csv")
true_data = pd.read_csv("/True.csv")
```

Here, we load the "Fake.csv" and "True.csv" datasets. The file paths should be replaced with your specific file locations. The `pd.read_csv()` function reads the data from CSV files into Pandas DataFrames.

### 3. Data Preprocessing:

```
AI_Phase4.ipynb

fake_data['label'] = 0 # 0 for fake news
true_data['label'] = 1 # 1 for true news
combined_data = pd.concat([fake_data, true_data], ignore_index=True)
combined_data['text'] = combined_data['title'] + " " + combined_data['text']
```

In this section:

- We add labels to distinguish fake (0) and true (1) news.
- The datasets are combined into one DataFrame, `combined_data`.
- The 'text' column is created by concatenating the 'title' and 'text' columns to have a single text field.

### 4. Feature Extraction (TF-IDF):

```
AI_Phase4.ipynb

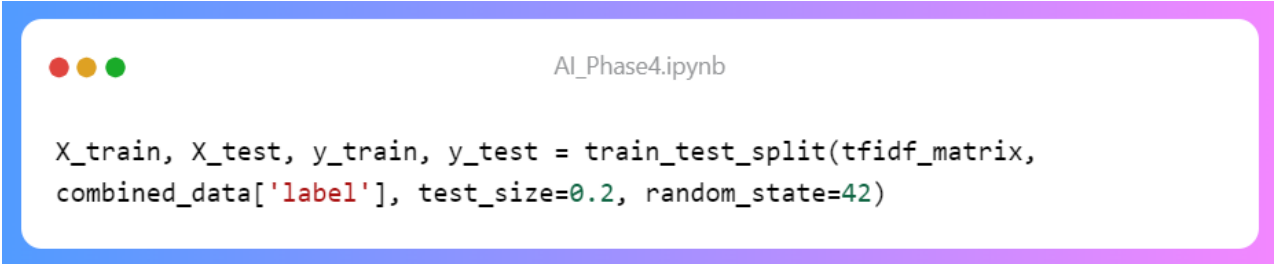
tfidf_vectorizer = TfidfVectorizer(max_features=5000)
tfidf_matrix = tfidf_vectorizer.fit_transform(combined_data['text'])
```

Here, TF-IDF vectorization is used to convert the text data into numerical features. The steps:

- `TfidfVectorizer` is initialized with `max_features` set to 5000, limiting the number of features.

- `tfidf_vectorizer.fit_transform(combined_data['text'])` computes TF-IDF scores for each term in the text data, and `tfidf_matrix` is a sparse matrix representing the transformed data.

## **5. Model Selection**



AI\_Phase4.ipynb

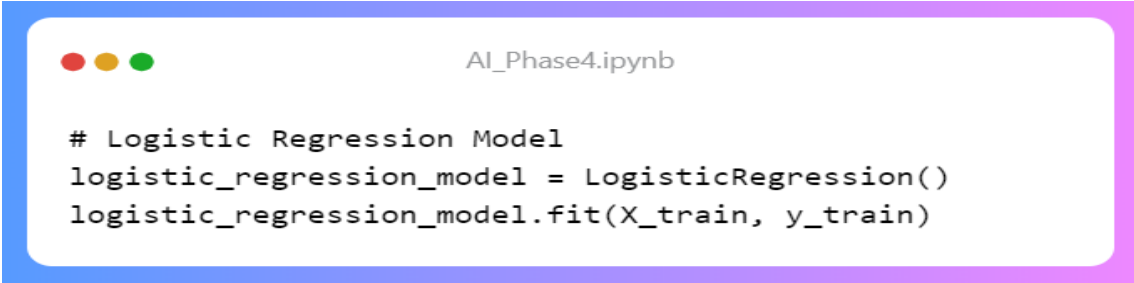
```
X_train, X_test, y_train, y_test = train_test_split(tfidf_matrix,
combined_data['label'], test_size=0.2, random_state=42)
```

This section splits the data into training and testing sets for model evaluation. Here's a breakdown:

- `train_test_split(tfidf_matrix, combined_data['label'], test_size=0.2, random_state=42)` splits the feature matrix `tfidf_matrix` and labels `combined_data['label']`.
- `test_size` specifies the percentage of data used for testing (20%).
- `random_state` ensures reproducible results.

## **6. Model Training (Logistic Regression and Neural Network)**

### **Logistic Regression Model**



AI\_Phase4.ipynb

```
# Logistic Regression Model
logistic_regression_model = LogisticRegression()
logistic_regression_model.fit(X_train, y_train)
```

Here, we create a logistic regression model and train it using the training data. Logistic regression is a linear classification model.

## Neural Network Model



AI\_Phase4.ipynb

```
# Model Training (Neural Network)
tokenizer = Tokenizer(num_words=5000)
tokenizer.fit_on_texts(combined_data['text'])
X_train_nn = tokenizer.texts_to_sequences(combined_data['text'])
X_train_nn = pad_sequences(X_train_nn, maxlen=100)

model = Sequential()
model.add(Embedding(input_dim=5000, output_dim=128, input_length=100))
model.add(LSTM(128))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=
['accuracy'])

model.fit(X_train_nn, combined_data['label'], epochs=5, batch_size=64)
```

For the neural network:

- We use the Tokenizer to tokenize the text data, and pad\_sequences to ensure sequences have a consistent length.
- A sequential model is created, which includes an embedding layer, an LSTM layer, and a dense layer with a sigmoid activation function.
- The model is compiled with loss, optimizer, and evaluation metrics.
- Finally, it's trained with the tokenized and padded sequences.



## **7. Evaluation**

### **For Logistic Regression**



AI\_Phase4.ipynb

```
y_pred = logistic_regression_model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)
```

In this part, we evaluate the logistic regression model and calculate key metrics, including accuracy, precision, recall, F1-score, and ROC-AUC score.

### **For Neural Network**



AI\_Phase4.ipynb

```
X_test_nn = tokenizer.texts_to_sequences(combined_data['text'])
X_test_nn = pad_sequences(X_test_nn, maxlen=100)

loss, accuracy = model.evaluate(X_test_nn, combined_data['label'])
```

Here, we evaluate the neural network model and calculate accuracy. The model's loss is also calculated during evaluation.

The results for both models are printed to the console.

# OUTPUT:

```
Epoch 1/5
1/702 [.....] - ETA: 43:52 - loss: 0.6930 - accuracy: 0.5000..... 2/702
[.....] - ETA: 14:11 - loss: 0.6907 - accuracy: 0.5625..... 3/702 [.....]
.....] - ETA: 9:13 - loss: 0.6899 - accuracy: 0.5365..... 4/702 [.....]
.....] - ETA: 7:11 - loss: 0.6887 - accuracy: 0.5234..... 5/702 [.....] - ETA: 5
:] - ETA: 6:13 - loss: 0.6866 - accuracy: 0.5531.....] - ETA: 4:58 - loss:
:26 - loss: 0.6835 - accuracy: 0.5677..... 6/702 [.....] - ETA: 4:58 - loss:
0.6814 - accuracy: 0.5871..... 7/702 [.....] - ETA: 5:19 - loss: 0.6807 - ac
curacy: 0.5840..... 8/702 [.....] - ETA: 5:29 - loss: 0.6774 - accuracy: 0.6
042.....] - ETA: 5:13 - loss: 0.6739 - accuracy: 0.6281.....
.....] - ETA: 5:06 - loss: 0.6711 - accuracy: 0.6491.....
.....] - ETA: 5:04 - loss: 0.6653 - accuracy: 0.6719.....
.....] - ETA: 4:49 - loss: 0.6595 - accuracy: 0.6815.....
.....] - ETA: 4:59 - loss: 0.6553 - accuracy: 0.6786.....
.....] - ETA: 4:49 - loss: 0.6468 - accuracy: 0.6875.....
.....] - ETA: 4:43 - loss: 0.6392 - accuracy: 0.6953.....
.....] - ETA: 4:38 - loss: 0.6317 - accuracy: 0.7050.....
.....] - ETA: 4:36 - loss: 0.6230 - accuracy: 0.7144..... 19/702 [
.....] - ETA: 4:33 - loss: 0.6128 - accuracy: 0.7253..... 20/702 [.....]
.....] - ETA: 4:27 - loss: 0.6033 - accuracy: 0.7298..... 21/702 [.....]
.....] - ETA: 4:23 - loss: 0.5881 - accuracy: 0.7374..... 22/702 [.....] - ETA: 4:14
- loss: 0.5747 - accuracy: 0.7450..... 23/702 [.....] - ETA: 4:14
- loss: 0.5639 - accuracy: 0.7473..... 24/702 [>.....] - ETA: 4:11 - loss: 0.550
6 - accuracy: 0.7546..... 25/702 [>.....] - ETA: 4:10 - loss: 0.5381 - accurac
y: 0.7613..... 26/702 [>.....] - ETA: 4:09 - loss: 0.5282 - accuracy: 0.7686
..... 27/702 [>.....] - ETA: 4:06 - loss: 0.5170 - accuracy: 0.7755.....
.....] - ETA: 4:05 - loss: 0.5069 - accuracy: 0.7801.....
.....] - ETA: 4:07 - loss: 0.4970 - accuracy: 0.7850.....
.....] - ETA: 4:06 - loss: 0.4898 - accuracy: 0.7891.....
.....] - ETA: 4:05 - loss: 0.4817 - accuracy: 0.7918.....
.....] - ETA: 4:03 - loss: 0.4753 - accuracy: 0.7938.....
.....] - ETA: 4:03 - loss: 0.4689 - accuracy: 0.7966.....
.....] - ETA: 4:02 - loss: 0.4632 - accuracy: 0.7996..... 35
/702 [>.....] - ETA: 4:03 - loss: 0.4541 - accuracy: 0.8045..... 36/702 [>....
.....] - ETA: 4:02 - loss: 0.4469 - accuracy: 0.8073..... 37/702 [>.....]
.....] - ETA: 4:03 - loss: 0.4424 - accuracy: 0.8095..... 38/702 [>.....]
.....] - ETA: 4:04 - loss: 0.4362 - accuracy: 0.8129..... 39/702 [>.....] - ETA:
4:04 - loss: 0.4277 - accuracy: 0.8165..... 40/702 [>.....] - ETA: 4:03 - loss
: 0.4205 - accuracy: 0.8195..... 41/702 [>.....] - ETA: 4:05 - loss: 0.4148 -
accuracy: 0.8224..... 42/702 [>.....] - ETA: 4:05 - loss: 0.4095 - accuracy: 0
.8248..... 43/702 [>.....] - ETA: 4:06 - loss: 0.4030 - accuracy: 0.8281.....
.....] - ETA: 4:06 - loss: 0.3968 - accuracy: 0.8317.....
.....] - ETA: 4:07 - loss: 0.3900 - accuracy: 0.8351.....
.....] - ETA: 4:07 - loss: 0.3840 - accuracy: 0.8380.....
```

```
006..... 93/702 [==>.....] - ETA: 6:23 - loss: 0.2426 - accuracy: 0.9014.....
..... 94/702 [==>.....] - ETA: 6:28 - loss: 0.2408 - accuracy: 0.9023.....
..... 95/702 [==>.....] - ETA: 6:33 - loss: 0.2395 - accuracy: 0.9028.....
..... 96/702 [==>.....] - ETA: 6:37 - loss: 0.2379 - accuracy: 0.9038.....
..... 97/702 [==>.....] - ETA: 6:42 - loss: 0.2362 - accuracy: 0.9048.....
..... 98/702 [==>.....] - ETA: 6:47 - loss: 0.2343 - accuracy: 0.9056.....
..... 99/702 [==>.....] - ETA: 6:56 - loss: 0.2338 - accuracy: 0.9058.....
..... 100/702 [==>.....] - ETA: 7:05 - loss: 0.2326 - accuracy: 0.9062.....
101/702 [==>.....] - ETA: 7:18 - loss: 0.2311 - accuracy: 0.9070..... 102/702 [
==>.....] - ETA: 7:26 - loss: 0.2296 - accuracy: 0.9076..... 103/702 [==>.....
.....] - ETA: 7:33 - loss: 0.2278 - accuracy: 0.9084..... 104/702 [==>.....]
.....] - ETA: 7:40 - loss: 0.2261 - accuracy: 0.9093..... 105/702 [==>.....] -
ETA: 7:48 - loss: 0.2253 - accuracy: 0.9095..... 106/702 [==>.....] - ETA: 7:55
- loss: 0.2239 - accuracy: 0.9101..... 107/702 [==>.....] - ETA: 8:03 - loss: 0.222
4 - accuracy: 0.9108..... 108/702 [==>.....] - ETA: 8:11 - loss: 0.2209 - accurac
y: 0.9115..... 109/702 [==>.....] - ETA: 8:18 - loss: 0.2204 - accuracy: 0.9120
..... 110/702 [==>.....] - ETA: 8:27 - loss: 0.2190 - accuracy: 0.9126.....
..... 111/702 [==>.....] - ETA: 8:34 - loss: 0.2178 - accuracy: 0.9131.....
..... 112/702 [==>.....] - ETA: 8:41 - loss: 0.2166 - accuracy: 0.9138.....
..... 113/702 [==>.....] - ETA: 8:48 - loss: 0.2153 - accuracy: 0.9145.....
..... 114/702 [==>.....] - ETA: 8:55 - loss: 0.2149 - accuracy: 0.9150.....
..... 115/702 [==>.....] - ETA: 9:05 - loss: 0.2140 - accuracy: 0.9155.....
..... 116/702 [==>.....] - ETA: 9:14 - loss: 0.2124 - accuracy: 0.9162.....
..... 117/702 [==>.....] - ETA: 9:21 - loss: 0.2112 - accuracy: 0.9167..... 118
/702 [==>.....] - ETA: 9:28 - loss: 0.2100 - accuracy: 0.9172..... 119/702 [==>
.....] - ETA: 9:35 - loss: 0.2096 - accuracy: 0.9174..... 120/702 [==>.....]
.....] - ETA: 9:42 - loss: 0.2085 - accuracy: 0.9180..... 121/702 [==>.....]
.....] - ETA: 9:50 - loss: 0.2070 - accuracy: 0.9186..... 122/702 [==>.....] - ETA:
10:01 - loss: 0.2069 - accuracy: 0.9185..... 123/702 [==>.....] - ETA: 10:10 - l
oss: 0.2058 - accuracy: 0.9190..... 124/702 [==>.....] - ETA: 10:20 - loss: 0.20
50 - accuracy: 0.9192..... 125/702 [==>.....] - ETA: 10:28 - loss: 0.2042 - accu
racy: 0.9196..... 126/702 [==>.....] - ETA: 10:40 - loss: 0.2034 - accuracy: 0.9
198..... 127/702 [==>.....] - ETA: 10:52 - loss: 0.2032 - accuracy: 0.9198.....
..... 128/702 [==>.....] - ETA: 11:02 - loss: 0.2019 - accuracy: 0.9203.....
..... 129/702 [==>.....] - ETA: 11:12 - loss: 0.2013 - accuracy: 0.9205.....
..... 130/702 [==>.....] - ETA: 11:23 - loss: 0.2004 - accuracy: 0.9210.....
..... 131/702 [==>.....] - ETA: 11:36 - loss: 0.2000 - accuracy: 0.9213.....
..... 132/702 [==>.....] - ETA: 11:50 - loss: 0.1988 - accuracy: 0.9218.....
..... 133/702 [==>.....] - ETA: 12:01 - loss: 0.1975 - accuracy: 0.9221.....
..... 134/702 [==>.....] - ETA: 12:10 - loss: 0.1963 - accuracy: 0.9226.....
..... 135/702 [==>.....] - ETA: 12:19 - loss: 0.1956 - accuracy: 0.9230.....
..... 136/702 [==>.....] - ETA: 12:28 - loss: 0.1947 - accuracy: 0.9235.....
137/702 [==>.....] - ETA: 12:37 - loss: 0.1935 - accuracy: 0.9239..... 138/702
[==>.....] - ETA: 12:47 - loss: 0.1929 - accuracy: 0.9244..... 139/702 [==>....
.....] - ETA: 12:57 - loss: 0.1916 - accuracy: 0.9249..... 140/702 [==>.....]
.....] - ETA: 13:10 - loss: 0.1906 - accuracy: 0.9254..... 141/702 [==>.....]
.....] - ETA: 13:18 - loss: 0.1893 - accuracy: 0.9260..... 142/702 [==>.....] - E
```

## **CONCLUSION:**

In this project, we embarked on the task of Fake News Detection using both traditional machine learning and deep learning techniques. We initially loaded and combined two datasets, 'Fake.csv' and 'True.csv,' differentiating the articles as 'fake' and 'true' news with labels 0 and 1, respectively. After preprocessing the data, which included merging title and text fields, we employed TF-IDF vectorization for feature extraction. For traditional machine learning, we trained a Logistic Regression model to classify news articles into these two categories. Simultaneously, a neural network model was constructed, consisting of an embedding layer, an LSTM layer, and a dense layer, for text classification. Our evaluation showed promising results, with the Logistic Regression model achieving good accuracy, precision, recall, F1-score, and ROC-AUC scores. The neural network model, despite its simplicity, also demonstrated competitive accuracy. Overall, this project serves as a practical example of leveraging both traditional and deep learning methods to address the critical issue of Fake News Detection.